

Intro to R

Data Cleaning

Data Cleaning

In general, data cleaning is a process of investigating your data for inaccuracies, or recoding it in a way that makes it more manageable.

 **MOST IMPORTANT RULE - LOOK  AT YOUR DATA! **

Dealing with Missing Data

Missing data types

One of the most important aspects of data cleaning is missing values.

Types of “missing” data:

- `NA` - general missing data
- `NaN` - stands for “**N**ot **a** **N**umber”, happens when you do $0/0$.
- `Inf` and `-Inf` - Infinity, happens when you take a positive number (or negative number) by 0.

Finding Missing data

Each missing data type has a function that returns `TRUE` if the data is missing:

- `NA - is.na`
- `NaN - is.nan`
- `Inf and -Inf - is.infinite`

Useful checking functions

- `is.na` - is TRUE if the data is FALSE otherwise
- `!` - negation (NOT)
 - if `is.na(x)` is TRUE, then `!is.na(x)` is FALSE
- `any` will be TRUE if ANY are true
 - `any(is.na(x))` - do we have any NA's in x?

```
A = c(1, 2, 4, NA)
B = c(1, 2, 3, 4)
any(is.na(A)) # are there any NAs - YES/TRUE
```

```
[1] TRUE
```

```
any(is.na(B)) # are there any NAs- NO/FALSE
```

```
[1] FALSE
```

naniar

Sometimes you need to look at lots of data though... the [naniar package](#) is a good option.

The `pct_complete()` function shows the percentage that is complete for a given data object.

```
#install.packages("naniar")  
library(naniar)  
x = c(0, NA, 2, 3, 4, -0.5, 0.2)  
naniar::pct_complete(x)
```

```
[1] 85.71429
```

Air quality data

The airquality dataset comes with R about air quality in New York in 1973.

```
?airquality # use this to find out more about the data
airqual <-tibble(airquality)
airqual
```

```
# A tibble: 153 × 6
   Ozone Solar.R Wind Temp Month Day
  <int>   <int> <dbl> <int> <int> <int>
1     41     190   7.4    67     5     1
2     36     118    8     72     5     2
3     12     149  12.6    74     5     3
4     18     313  11.5    62     5     4
5    NA      NA  14.3    56     5     5
6     28      NA  14.9    66     5     6
7     23     299   8.6    65     5     7
8     19      99  13.8    59     5     8
9      8      19  20.1    61     5     9
10    NA     194   8.6    69     5    10
# ... with 143 more rows
```


naniar: `pct_complete()`


```
pct_complete(airquality)
```

```
[1] 95.20697
```

Naniar plots

The `gg_miss_var()` function creates a nice plot about the number of missing values for each variable.


```
naniar::gg_miss_var(airqual)
```



Naniar plots

We can use the `facet` argument to make more plots about a specific variable.

```
naniar::gg_miss_var(airqual, facet = Month)
```



Missing Data Issues

Recall that mathematical operations with NA often result in NAs.

```
sum(c(1, 2, 3, NA))
```

```
[1] NA
```

```
mean(c(2, 4, NA))
```

```
[1] NA
```

```
median(c(1, 2, 3, NA))
```

```
[1] NA
```

Missing Data Issues

This is also true for logicals. This is a good thing. The NA data could be > 2 or not, we don't know, so R says there is no TRUE or FALSE, so that is missing.

```
x = c(0, NA, 2, 3, 4, -0.5, 0.2)
x > 2
```

```
[1] FALSE      NA FALSE    TRUE    TRUE FALSE  FALSE
```

filter() and missing data

Be careful with missing data using subsetting:

`filter()` removes missing values by default. To keep them need to add `is.na()`:

```
x # looks like the 1st and 3rd element should be TRUE
```

```
[1] 0.0 NA 2.0 3.0 4.0 -0.5 0.2
```

```
x %in% c(0, 2) # uh oh - not good!
```

```
[1] TRUE FALSE TRUE FALSE FALSE FALSE FALSE
```

```
x %in% c(0, 2) | is.na(x) # do this
```

```
[1] TRUE TRUE TRUE FALSE FALSE FALSE FALSE
```

filter() and missing data

```
df
```

```
# A tibble: 6 × 2
```

	Dog	Cat
	<dbl>	<dbl>
1	0	NA
2	NA	8
3	2	6
4	3	NA
5	1	2
6	1	NA

```
df %>% filter(Dog < 3)
```

```
# A tibble: 4 × 2
```

	Dog	Cat
	<dbl>	<dbl>
1	0	NA
2	2	6
3	1	2
4	1	NA

to filter out NAs for one variable

```
df %>% filter(!is.na(Dog))
```

```
# A tibble: 5 × 2
  Dog   Cat
<dbl> <dbl>
1     0   NA
2     2     6
3     3   NA
4     1     2
5     1   NA
```

`!NA` does not work as you might expect because you can't tell if something is not actually `NA` - R doesn't ever assume to know what the value of `NA` is

```
NA == NA
```

```
[1] NA
```

```
NA != NA
```

```
[1] NA
```


tidyr::drop_na

This function will drop rows with **any** missing data in **any** column.

```
df
```

```
# A tibble: 6 × 2
  Dog   Cat
  <dbl> <dbl>
1     0   NA
2    NA     8
3     2     6
4     3   NA
5     1     2
6     1   NA
```

```
drop_na(df)
```

```
# A tibble: 2 × 2
  Dog   Cat
  <dbl> <dbl>
1     2     6
2     1     2
```

Think about **NA**

Sometimes removing NA values leads to distorted math - be careful! Think about what your NA means for your data (are you sure?).

Is an NA for values so low they could not be reported? Or is it this and also if there was a different issue?

Think about **NA**

If it is something more like a zero then you might want it included in your data like a zero.

Example: - survey reports NA if student has never tried cigarettes - survey reports 0 if student has tried cigarettes but did not smoke that week

You might want to keep the NA values so that you know the original sample size.

Word of caution

Calculating percentages will give you a different result depending on your choice to include NA values.

```
red_blue
```

```
# A tibble: 3 × 2
  color col_count
  <chr>     <int>
1 blue         3
2 red          3
3 <NA>         3
```

```
red_blue %>% mutate(percent =
  col_count/sum(pull(red_blue, col_count)))
```

```
# A tibble: 3 × 3
  color col_count percent
  <chr>     <int>   <dbl>
1 blue         3    0.333
2 red          3    0.333
3 <NA>         3    0.333
```

Word of caution

```
red_blue %>% mutate(percent =  
  col_count/sum(pull(drop_na(red_blue), col_count)))
```

```
# A tibble: 3 × 3  
  color col_count percent  
  <chr>    <int>    <dbl>  
1 blue         3      0.5  
2 red          3      0.5  
3 <NA>         3      0.5
```

```
# Should you be dividing by 9 or 6? It depends on your data  
# Pay attention to your data and your NAs!
```

Check values

Check the values for your variables, are they what you expect?

`count()` is a great option because it gives tells you:

1. The unique values
2. the amount of these values

Check if rare values make sense

```
bike <- jhur::read_bike()
```

```
bike %>% count(subType)
```

```
# A tibble: 4 × 2
```

	subType	n
	<chr>	<int>
1	STCLN	1
2	STRALY	3
3	STRPRD	1623
4	<NA>	4

```
bike %>% pull(subType) %>% unique()
```

```
[1] NA      "STCLN" "STRALY" "STRPRD"
```

Lab Part 1

[lab part 1](#)

[Website](#)

Recoding Variables

Example of Recoding

Say we have some data about samples in a diet study:

```
data_diet
```

```
# A tibble: 12 × 4
  Diet   Gender Weight_start Weight_change
<chr> <chr>      <int>      <int>
1 A     Male      147         2
2 B     m        134         8
3 B     Other     209         7
4 A     F         127        -7
5 B     Female     239        15
6 B     M         230       -10
7 A     f         199        -8
8 B     O         156         9
9 B     Man       135        19
10 A    f         152         4
11 B    F         242        17
12 B    O         108         5
```

Oh dear...

This needs lots of recoding.

```
data_diet %>%  
  count(Gender, Diet)
```

```
# A tibble: 10 × 3  
  Gender Diet      n  
  <chr>  <chr> <int>  
1 f      A      2  
2 F      A      1  
3 F      B      1  
4 Female B      1  
5 m      B      1  
6 M      B      1  
7 Male   A      1  
8 Man    B      1  
9 O      B      2  
10 Other B      1
```

dplyr can help!

Using Excel to find all of the different ways `gender` has been coded, would be a matter of filtering and changing all by hand or using if statements. This can be hectic!

In `dplyr` you can use the `recode` function (need `mutate` here too!):

```
# General Format - this is not code!
{data_input} %>%
  mutate({variable_to_fix} = {Variable_fixing, {old_value} = {new_value},
                                     {another_old_value} = {new_value}})
```

```
data_diet %>%
  mutate(Gender = recode(Gender, M = "Male",
                           m = "Male",
                           Man = "Male",
                           O = "Other",
                           f = "Female",
                           F = "Female")) %>%
  count(Gender, Diet)
```

Or you can use `case_when()`.

The `case_when()` function of `dplyr` can help us to do this as well.

```
# General Format - this is not code!  
{data_input} %>%  
  mutate({variable_to_fix} = case_when{Variable_fixing}/condition/  
                                             ~{value_for_cond}))
```

Note that automatically values not reassigned explicitly by `case_when` will be NA.

Use of `case_when()`

```
data_diet %>%  
  mutate(Gender = case_when(Gender == "M" ~ "Male"))
```

```
# A tibble: 12 × 4  
  Diet Gender Weight_start Weight_change  
  <chr> <chr>      <int>      <int>  
1 A     <NA>      147         2  
2 B     <NA>      134         8  
3 B     <NA>      209         7  
4 A     <NA>      127        -7  
5 B     <NA>      239        15  
6 B     Male      230       -10  
7 A     <NA>      199        -8  
8 B     <NA>      156         9  
9 B     <NA>      135        19  
10 A    <NA>      152         4  
11 B    <NA>      242        17  
12 B    <NA>      108         5
```

More complicated case_when()

```
data_diet %>%  
  mutate(Gender = case_when(  
    Gender %in% c("M", "male", "Man", "m", "Male") ~ "Male",  
    Gender %in% c("F", "Female", "f", "female") ~ "Female",  
    Gender %in% c("O", "Other") ~ "Other"))
```

```
# A tibble: 12 × 4  
  Diet Gender Weight_start Weight_change  
  <chr> <chr>      <int>      <int>  
1 A     Male        147         2  
2 B     Male        134         8  
3 B     Other        209         7  
4 A     Female        127        -7  
5 B     Female        239        15  
6 B     Male         230       -10  
7 A     Female        199        -8  
8 B     Other        156         9  
9 B     Male         135        19  
10 A    Female        152         4  
11 B    Female        242        17  
12 B    Other         108         5
```

Another reason for `case_when()`

`case_when` can do very sophisticated comparisons

```
data_diet <- data_diet %>%  
  mutate(Effect = case_when(Weight_change > 0 ~ "Increase",  
                             Weight_change == 0 ~ "Same",  
                             Weight_change < 0 ~ "Decrease"))
```

`data_diet`

```
# A tibble: 12 × 5  
  Diet Gender Weight_start Weight_change Effect  
  <chr> <chr>      <int>      <int> <chr>  
1 A     Male      147         2 Increase  
2 B     m        134         8 Increase  
3 B     Other    209         7 Increase  
4 A     F        127        -7 Decrease  
5 B     Female   239        15 Increase  
6 B     M        230       -10 Decrease  
7 A     f        199        -8 Decrease  
8 B     O        156         9 Increase  
9 B     Man      135        19 Increase  
10 A    f        152         4 Increase  
11 B    F        242        17 Increase  
12 B    O        108         5 Increase
```

```
data_diet %>%  
  count(Diet, Effect)
```

What if our data looked like this?

```
diet_comb
```

```
# A tibble: 4 × 2
  change      n
  <chr>    <int>
1 A_Decrease    2
2 A_Increase    2
3 B_Decrease    1
4 B_Increase    7
```


Separating columns based on a separator

- From `tidyr`, you can split a data set into multiple columns:

```
diet_comb %>%  
  separate(change, into = c("Diet", "Change"))
```

```
# A tibble: 4 × 3  
  Diet Change      n  
  <chr> <chr>   <int>  
1 A     Decrease  2  
2 A     Increase  2  
3 B     Decrease  1  
4 B     Increase  7
```

Separating columns based on a separator

You can specify the separator with `sep`.

```
diet_comb %>%  
  separate(change, into = c("Diet", "Change"), sep = " ")
```

```
# A tibble: 4 × 3  
  Diet      Change      n  
  <chr>   <chr>   <int>  
1 A_diet Decrease     2  
2 A_diet Increase     2  
3 B_diet Decrease     1  
4 B_diet Increase     7
```

Uniting columns based on a separator

- From `tidyr`, you can unite:

```
df = tibble(id = rep(1:5, 3), visit = rep(1:3, each = 5))  
head(df, 4)
```

```
# A tibble: 4 × 2  
  id visit  
  <int> <int>  
1     1     1  
2     2     1  
3     3     1  
4     4     1
```

```
df_united <- df %>% unite(col = "unique_id", id, visit, sep = "_")  
head(df_united, 4)
```

```
# A tibble: 4 × 1  
  unique_id  
  <chr>  
1 1_1  
2 2_1  
3 3_1  
4 4_1
```

Strings functions

Splitting/Find/Replace and Regular Expressions

- R can do much more than find exact matches for a whole string!

The **stringr** package

The `stringr` package:

- Modifying or finding **part** or all of a character string
- We will not cover `grep` or `gsub` - base R functions
 - are used on forums for answers
- Almost all functions start with `str_*`

stringr

`str_detect`, and `str_replace` search for matches to argument pattern within each element of a character vector (not data frame or tibble!).

- `str_detect` - returns TRUE if pattern is found
- `str_replace` - replaces pattern with replacement

Download Salary FY2014 Data

From <https://data.baltimorecity.gov/City-Government/Baltimore-City-Employee-Salaries-FY2015/nsfe-bg53>, from <https://data.baltimorecity.gov/api/views/nsfe-bg53/rows.csv>

Read the CSV into R `sal`:

```
sal = jhur::read_salaries() # or
```

```
head(sal)
```

```
# A tibble: 6 × 7
```

	name <chr>	JobTitle <chr>	AgencyID <chr>	Agency <chr>	HireDate <chr>	AnnualSalary <chr>	GrossPay <chr>
1	Aaron, Patricia G	Faciliti...	A03031	OED-Em...	10/24/1...	\$55314.00	\$53626
2	Aaron, Petra L	ASSISTAN...	A29045	States...	09/25/2...	\$74000.00	\$73000
3	Abaineh, Yohannes T	EPIDEMIO...	A65026	HLTH-H...	07/23/2...	\$64500.00	\$64403
4	Abbene, Anthony M	POLICE O...	A99005	Police...	07/24/2...	\$46309.00	\$59620
5	Abbey, Emmanuel	CONTRACT...	A40001	M-R In...	05/01/2...	\$60060.00	\$54059
6	Abbott-Cole, Michelle	CONTRACT...	A90005	TRANS-...	11/28/2...	\$42702.00	\$20250

'Find' functions: finding values: **stringr**

```
Sal %>% filter(str_detect(name, "Rawlings"))
```

```
# A tibble: 3 × 7
```

	name	JobTitle	AgencyID	Agency	HireDate	AnnualSalary	GrossE
	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>
1	Rawlings, Ke...	EMERGENCY D...	A40302	M-R Info Te...	01/06/2...	\$48940.00	\$73356
2	Rawlings, Pa...	COMMUNITY A...	A04015	R&P-Recreat...	12/10/2...	\$19802.00	\$10443
3	Rawlings-Bl...	MAYOR	A01001	Mayors Offi...	12/07/1...	\$167449.00	\$16524

Lab Part 2

[lab part 2](#)

[Website](#)

Extra Slides

String Splitting

- `str_split(string, pattern)` - splits strings up - returns list!

```
library(stringr)
x <- c("I really like writing R code")
df = tibble(x = c("I really", "like writing", "R code programs"))
y <- unlist(str_split(x, " "))
y
```

```
[1] "I"          "really"    "like"      "writing"   "R"         "code"
```

```
length(y)
```

```
[1] 6
```

A bit on Regular Expressions

- <http://www.regular-expressions.info/reference.html>
- They can use to match a large number of strings in one statement
- `.` matches any single character
- `*` means repeat as many (even if 0) more times the last character
- `?` makes the last thing optional
- `^` matches start of vector `^a` - starts with "a"
- `$` matches end of vector `b$` - ends with "b"

Let's look at modifiers for **stringr**

?modifiers

- `fixed` - match everything exactly
- `ignore_case` is an option to not have to use `tolower`

Using a fixed expression

One example case is when you want to split on a period ".". In regular expressions . means **ANY** character, so we need to specify that we want R to interpret "." as simply a period.

```
str_split("I.like.strings", ".")
```

```
[[1]]  
[1] "" "" "" "" "" "" "" "" "" "" "" "" "" "" ""
```

```
str_split("I.like.strings", fixed("."))
```

```
[[1]]  
[1] "I"      "like"   "strings"
```

```
str_split("I.like.strings", "\\.")
```

```
[[1]]  
[1] "I"      "like"   "strings"
```

Showing difference in `str_replace` and `str_replace_all`

`str_replace` replaces only the first instance.

```
head(Sal$Name, 2)
```

```
Warning: Unknown or uninitialised column: `Name`.
```

```
NULL
```

```
head(str_replace(Sal$name, "a", "j"), 2)
```

```
[1] "Ajron, Patricia G" "Ajron, Petra L"
```

`str_replace` replaces all instances.

```
head(str_replace_all(Sal$name, "a", "j"), 2)
```

```
[1] "Ajron, Pjtricij G" "Ajron, Petrj L"
```


Pasting strings with `paste` and `paste0`

Paste can be very useful for joining vectors together:

```
paste("Visit", 1:5, sep = "_")
```

```
[1] "Visit_1" "Visit_2" "Visit_3" "Visit_4" "Visit_5"
```

```
paste("Visit", 1:5, sep = "_", collapse = "_")
```

```
[1] "Visit_1_Visit_2_Visit_3_Visit_4_Visit_5"
```

```
# and paste0 can be even simpler see ?paste0  
paste0("Visit", 1:5) # no space!
```

```
[1] "Visit1" "Visit2" "Visit3" "Visit4" "Visit5"
```

!- # Before Cleaning - Subsetting with Brackets ->

->

-> -> ->

Using Regular Expressions

- Look for any name that starts with:
 - Payne at the beginning,
 - Leonard and then an S
 - Spence then capital C

```
head(str_subset( Sal$name, "^Payne.*"), 3)
```

```
[1] "Payne El,Boaz L"      "Payne El,Jackie"  
[3] "Payne Johnson,Nickole A"
```

```
head(str_subset( Sal$name, "Leonard.?S"))
```

```
[1] "Payne,Leonard S"      "Szumlanski,Leonard S"
```

```
head(str_subset( Sal$name, "Spence.*C.*"))
```

```
[1] "Spencer,Charles A"  "Spencer,Clarence W" "Spencer,Michael C"
```

Comparison of **stringr** to base R -
not covered

Splitting Strings

Substringing

stringr

- `str_split(string, pattern)` - splits strings up - returns list!

Splitting String:

In `stringr`, `str_split` splits a vector on a string into a list

```
x <- c("I really", "like writing", "R code programs")  
y <- stringr::str_split(x, pattern = " ") # returns a list  
y
```

```
[[1]]  
[1] "I"      "really"
```

```
[[2]]  
[1] "like"   "writing"
```

```
[[3]]  
[1] "R"      "code"   "programs"
```

str_extract

`str_extract` extracts matched strings - `\\d` searches for DIGITS/numbers

```
head(Sal$AgencyID)
```

```
[1] "A03031" "A29045" "A65026" "A99005" "A40001" "A90005"
```

```
head(str_extract(Sal$AgencyID, "\\d"))
```

```
[1] "0" "2" "6" "9" "4" "9"
```

'Find' functions: stringr compared to base R

Base R does not use these functions. Here is a "translator" of the `stringr` function to base R functions

- `str_detect` - similar to `grepl` (return logical)
- `grep(value = FALSE)` is similar to `which(str_detect())`
- `str_subset` - similar to `grep(value = TRUE)` - return value of matched
- `str_replace` - similar to `sub` - replace one time
- `str_replace_all` - similar to `gsub` - replace many times

Important Comparisons

Base R:

- Argument order is `(pattern, x)`
- Uses option `(fixed = TRUE)`

`stringr`

- Argument order is `(string, pattern)` aka `(x, pattern)`
- Uses function `fixed(pattern)`

'Find' functions: Finding Indices

These are the indices where the pattern match occurs:

```
grep("Rawlings", Sal$Name)
```

Warning: Unknown or uninitialised column: `Name`.

```
integer(0)
```

```
which(grepl("Rawlings", Sal$Name))
```

Warning: Unknown or uninitialised column: `Name`.

```
integer(0)
```

```
which(str_detect(Sal$Name, "Rawlings"))
```

Warning: Unknown or uninitialised column: `Name`.

```
integer(0)
```

'Find' functions: Finding Logicals

These are the indices where the pattern match occurs:

```
head(grepl("Rawlings", Sal$Name))
```

Warning: Unknown or uninitialised column: `Name`.

```
logical(0)
```

```
head(str_detect(Sal$Name, "Rawlings"))
```

Warning: Unknown or uninitialised column: `Name`.

```
logical(0)
```

'Find' functions: finding values, base R

```
grep("Rawlings", Sal$Name, value=TRUE)
```

Warning: Unknown or uninitialised column: `Name`.

```
character(0)
```

```
Sal[grep("Rawlings", Sal$Name), ]
```

Warning: Unknown or uninitialised column: `Name`.

```
# A tibble: 0 × 7
```

```
# ... with 7 variables: name <chr>, JobTitle <chr>, AgencyID <chr>, Agency <chr>
```

```
#   HireDate <chr>, AnnualSalary <chr>, GrossPay <chr>
```

Showing difference in `str_extract`

`str_extract` extracts just the matched string

```
ss = str_extract(Sal$Name, "Rawling")
```

Warning: Unknown or uninitialised column: `Name`.

```
head(ss)
```

```
character(0)
```

```
ss[ !is.na(ss) ]
```

```
character(0)
```

Showing difference in `str_extract` and `str_extract_all`

`str_extract_all` extracts all the matched strings

```
head(str_extract(Sal$AgencyID, "\\d"))
```

```
[1] "0" "2" "6" "9" "4" "9"
```

```
head(str_extract_all(Sal$AgencyID, "\\d"), 2)
```

```
[[1]]
```

```
[1] "0" "3" "0" "3" "1"
```

```
[[2]]
```

```
[1] "2" "9" "0" "4" "5"
```

Using Regular Expressions

- Look for any name that starts with:
 - Payne at the beginning,
 - Leonard and then an S
 - Spence then capital C

```
head(grep("^Payne.*", x = Sal$name, value = TRUE), 3)
```

```
[1] "Payne El,Boaz L"      "Payne El,Jackie"  
[3] "Payne Johnson,Nickole A"
```

```
head(grep("Leonard.?S", x = Sal$name, value = TRUE))
```

```
[1] "Payne,Leonard S"      "Szumlanski,Leonard S"
```

```
head(grep("Spence.*C.*", x = Sal$name, value = TRUE))
```

```
[1] "Spencer,Charles A"  "Spencer,Clarence W" "Spencer,Michael C"
```

Using Regular Expressions: **stringr**

```
head(str_subset( Sal$name, "^Payne.*"), 3)
```

```
[1] "Payne El,Boaz L"      "Payne El,Jackie"  
[3] "Payne Johnson,Nickole A"
```

```
head(str_subset( Sal$name, "Leonard.?S"))
```

```
[1] "Payne,Leonard S"      "Szumlanski,Leonard S"
```

```
head(str_subset( Sal$name, "Spence.*C.*"))
```

```
[1] "Spencer,Charles A"  "Spencer,Clarence W" "Spencer,Michael C"
```


Replace

Let's say we wanted to sort the data set by Annual Salary:

```
class(Sal$AnnualSalary)
```

```
[1] "character"
```

```
sort(c("1", "2", "10")) # not sort correctly (order simply ranks the data)
```

```
[1] "1"  "10" "2"
```

```
order(c("1", "2", "10"))
```

```
[1] 1 3 2
```

Replace

So we must change the annual pay into a numeric:

```
head(Sal$AnnualSalary, 4)
```

```
[1] "$55314.00" "$74000.00" "$64500.00" "$46309.00"
```

```
head(as.numeric(Sal$AnnualSalary), 4)
```

```
Warning in head(as.numeric(Sal$AnnualSalary), 4): NAs introduced by coercion
```

```
[1] NA NA NA NA
```

R didn't like the \$ so it thought turned them all to NA.

`sub()` and `gsub()` can do the replacing part in base R.

Replacing and subbing

Now we can replace the \$ with nothing (used `fixed=TRUE` because \$ means ending):

```
Sal$AnnualSalary <- as.numeric(gsub(pattern = "$", replacement="",  
                                   Sal$AnnualSalary, fixed=TRUE))  
Sal <- Sal[order(Sal$AnnualSalary, decreasing=TRUE), ]  
Sal[1:5, c("name", "AnnualSalary", "JobTitle")]
```

```
# A tibble: 5 × 3  
  name AnnualSalary JobTitle  
  <chr>      <dbl> <chr>  
1 Mosby, Marilyn J      238772 STATE'S ATTORNEY  
2 Batts, Anthony W      211785 Police Commissioner  
3 Wen, Leana           200000 Executive Director III  
4 Raymond, Henry J      192500 Executive Director III  
5 Swift, Michael        187200 CONTRACT SERV SPEC II
```

Replacing and subbing: **stringr**

We can do the same thing (with 2 piping operations!) in dplyr

```
dplyr_sal = Sal
dplyr_sal = dplyr_sal %>% mutate(
  AnnualSalary = AnnualSalary %>%
    str_replace(
      fixed("$"),
      "") %>%
    as.numeric() %>%
    arrange(desc(AnnualSalary))
check_Sal = Sal
rownames(check_Sal) = NULL
all.equal(check_Sal, dplyr_sal)
```

```
[1] TRUE
```

Website

Website

Extra slides

Creating Two-way Tables

A two-way table. If you pass in 2 vectors, `table` creates a 2-dimensional table.

```
tab <- table(c(0, 1, 2, 3, 2, 3, 3, 2, 2, 3),  
             c(0, 1, 2, 3, 2, 3, 3, 4, 4, 3),  
             useNA = "always")  
tab
```

	0	1	2	3	4	<NA>
0	1	0	0	0	0	0
1	0	1	0	0	0	0
2	0	0	2	0	2	0
3	0	0	0	4	0	0
<NA>	0	0	0	0	0	0

Creating Two-way Tables

```
tab_df = tibble(x = c(0, 1, 2, 3, 2, 3, 3, 2, 2, 3),  
                 y = c(0, 1, 2, 3, 2, 3, 3, 4, 4, 3))  
tab_df %>% count(x, y)
```

```
# A tibble: 5 × 3  
      x     y     n  
  <dbl> <dbl> <int>  
1     0     0     1  
2     1     1     1  
3     2     2     2  
4     2     4     2  
5     3     3     4
```


Creating Two-way Tables

```
tab_df %>%  
  count(x, y) %>%  
  group_by(x) %>% mutate(pct_x = n / sum(n))
```

```
# A tibble: 5 × 4  
# Groups:   x [4]  
      x     y     n pct_x  
  <dbl> <dbl> <int> <dbl>  
1     0     0     1     1  
2     1     1     1     1  
3     2     2     2    0.5  
4     2     4     2    0.5  
5     3     3     4     1
```

Creating Two-way Tables

```
library(scales)
tab_df %>%
  count(x, y) %>%
  group_by(x) %>% mutate(pct_x = percent(n / sum(n)))
```

```
# A tibble: 5 × 4
# Groups:   x [4]
      x     y     n pct_x
  <dbl> <dbl> <int> <chr>
1     0     0     1 100%
2     1     1     1 100%
3     2     2     2  50%
4     2     4     2  50%
5     3     3     4 100%
```