Intro to R

Functions

Writing your own functions

So far we've seen many functions, like `c()`, `class()`, `filter()`, `dim()` …

Why create your own functions?

- Cut down on repetitive code (easier to fix things!)
- Organize code into manageable chunks
- Avoid running code unintentionally
- Use names that make sense to you

Writing your own functions

Here we will write a function that multiplies some number (x) by 2:

```r
times_2 <- function(x) x * 2
```

When you run the line of code above, you make it ready to use (no output yet!). Let's test it!

```r
times_2(x = 10)
```

```
[1] 20
```

Writing your own functions: { }

Adding the curly brackets - { } - allows you to use functions spanning multiple lines:

```
times_2 <- function(x) {
  x * 2
}
times_2(x = 10)
```

```
[1] 20
```

Writing your own functions: `return`

If we want something specific for the function's output, we use `return()`:

```r
times_2 <- function(x) {
  output <- x * 2
  return(output)
}
times_2(x = 10)
```

```
[1] 20
```

Writing your own functions

Review: The syntax for a function is:

```
functionName = function(inputs) {
 <function body>
return(value)
}
```

Writing your own functions: multiple inputs

Functions can take multiple inputs.

```
times_2_plus_y <- function(x, y) x * 2 + y
times_2_plus_y(x = 10, y = 3)
```

```
[1] 23
```

Writing your own functions: defaults

Functions can have "default" arguments. This lets us use the function without using an argument later.

```
times_2_plus_y <- function(x = 10, y = 3) x * 2 + y
times_2_plus_y()
```

```
[1] 23
```

Writing another simple function

Let's write a function, `sqdif`, that:

1. takes two numbers `x` and `y` with default values of 2 and 3.

2. takes the difference

3. squares this difference

4. then returns the final value

Writing another simple function

```r
sqdif <- function(x = 2, y = 3) (x - y) ^ 2

sqdif()
```

```
[1] 1
```

```r
sqdif(x = 10, y = 5)
```

```
[1] 25
```

```r
sqdif(10, 5)
```

```
[1] 25
```

Writing your own functions: characters

Functions can have any kind of input. Here is a function with characters `return()`:

```
loud <- function(word) {
  output <- rep( toupper(word), 5)
  return(output)
}
loud("hooray!")
```

```
[1] "HOORAY!" "HOORAY!" "HOORAY!" "HOORAY!" "HOORAY!"
```

Functions for tibbles

We can use `filter(row_number()==n)` to extract a row of a tibble.

```r
cars <- read_kaggle()

get_row <- function(dat, row) dat %>% filter(row_number()==row)
```

```r
get_row(dat = cars, row = 10)
```

```
# A tibble: 1 x 10
  RefId IsBadBuy PurchDate Auction VehYear VehicleAge Make  Model Trim  SubMo
  <dbl>    <dbl> <chr>     <chr>     <dbl>      <dbl> <chr> <chr> <chr> <chr>
1    10        0 12/7/2009 ADESA      2007          2 FORD  FIVE… SEL   4D SED
```

Functions for tibbles

`select(n)` will choose column `n`.

```r
get_index <- function(dat, row, col){
 dat %>% filter(row_number()==row) %>% select(col)
}

get_index(dat = cars, row = 10, col = 8)
```

```
# A tibble: 1 x 1
  Model
  <chr>
1 FIVE HUNDRED
```

Functions for tibbles

Including default values for arguments:

```r
get_top <- function(dat, row = 1, col = 1){
 dat %>% filter(row_number()==row)  %>% select(col)
}

get_top(dat = cars)

# A tibble: 1 x 1
  RefId
  <dbl>
1     1
```

Using your custom functions: `sapply()`

You can also "apply" functions easily with `sapply()`.

These functions take the form:

```
sapply(<a vector or list>, some_function)
```

🚨 There are no parentheses on the functions! 🚨

Using your custom functions: `sapply()`

```
sapply(cars, class)
```

```
                      RefId                          IsBadBuy
                  "numeric"                         "numeric"
                 PurchDate                           Auction
                "character"                       "character"
                   VehYear                        VehicleAge
                  "numeric"                         "numeric"
                      Make                             Model
                "character"                       "character"
                      Trim                          SubModel
                "character"                       "character"
                     Color                      Transmission
                "character"                       "character"
               WheelTypeID                         WheelType
                "character"                       "character"
                    VehOdo                       Nationality
                  "numeric"                       "character"
                      Size             TopThreeAmericanName
                "character"                       "character"
MMRAcquisitionAuctionAveragePrice    MMRAcquisitionAuctionCleanPrice
                "character"                       "character"
 MMRAcquisitionRetailAveragePrice     MMRAcquisitonRetailCleanPrice
                "character"                       "character"
   MMRCurrentAuctionAveragePrice       MMRCurrentAuctionCleanPrice
                "character"                       "character"
    MMRCurrentRetailAveragePrice        MMRCurrentRetailCleanPrice
                "character"                       "character"
                  PRIMEUNIT                          AUCGUART
                "character"                       "character"
```

Using your custom functions "on the fly" to iterate

```r
sapply(pull(cars, VehOdo), function(x) x / 1000)
```

```
  [1]  89.046  93.593  73.807  65.617  69.367  81.054  65.328  65.805  49.92
 [10]  84.872  80.080  75.419  79.315  71.254  74.722  72.132  80.736  75.15
 [19]  65.925  84.498  54.586  66.536  98.130  59.789  65.663  52.106  88.95
 [28]  76.173  65.393  80.064  77.694  56.300  78.241  57.723  78.434  82.94
 [37]  76.304  55.711  76.586  65.078  65.403  86.889  68.990  80.949  52.77
 [46]  72.191  59.858  79.576  73.291  50.227  82.146  58.024  40.919  87.64
 [55]  80.968  50.308  80.795  62.239  87.008  64.060  77.677  58.888  63.55
 [64]  90.026  89.705  64.511  75.513  80.608  95.558  35.796  83.501  70.14
 [73]  76.052  72.479  84.542  61.081  86.483  43.898  57.338  59.425  79.95
 [82]  78.559  48.386  80.117  65.795  51.145  88.366  55.909  86.702  81.42
 [91]  65.379  74.954  49.328  73.810  43.412  78.412  74.026  64.822  80.49
[100]  85.003  65.711  56.064  62.230  62.190  67.426  75.806  88.991  89.84
[109]  81.338  80.077  77.233  66.681  82.526  81.930  74.131  72.417  64.11
[118]  71.423  64.650  85.388  95.443  69.337  46.563  84.905  71.062  80.99
[127]  66.545  67.785  71.952  70.741  94.318  69.440  54.268  59.072  86.02
[136]  64.677  68.874  64.554  73.988  23.881  50.532  60.554  91.558  63.37
[145]  59.391  44.367  44.515  83.238  92.532  68.165  87.775  86.414  36.14
[154]  80.788  93.346  73.963  68.183  64.839  75.484  59.287  63.151  46.69
[163]  58.897  65.363  75.237  85.042  87.701  92.816  97.221  73.726  47.55
[172]  63.079  64.064  88.027  82.164  84.763  52.113  49.893  92.782  46.00
[181]  62.990  78.992  64.458  60.522  73.725  71.214  60.530  66.695  89.03
[190]  36.425  58.823  72.592  79.015  88.667  58.499  95.025  50.644  88.83
[199]  68.040  58.384  79.284  80.906  94.011  86.875  61.319  79.333  92.89
[208]  59.801  75.108  67.696  50.385  58.450  75.070  73.870  43.535  55.68
[217]  58.681  62.795  77.178  69.430  86.466  69.480  86.344  79.030  94.50
[226]  77.786  67.430  73.409  60.201  49.296  61.315  84.454  85.262  72.99
[235]  89.769  68.550  62.955  75.101  88.173  52.791  85.334  69.030  64.55
[244]  66.433  77.050  84.388  52.866  45.680  63.496  78.593  93.395  53.46
```

Using your custom functions: `sapply()`.

```r
cars_dbl <- cars %>% select(Make, Model, where(is.double))


Odo_updated <- sapply(pull(cars_dbl, VehOdo), times_2_plus_y)


cars_dbl %>%
  mutate(Odo_2_y = Odo_updated)
```

```
# A tibble: 72,983 x 13
    Make    Model    RefId IsBadBuy VehYear VehicleAge VehOdo BYRNO VNZIP1 VehBCost
    <chr>   <chr>    <dbl>    <dbl>   <dbl>      <dbl>  <dbl> <dbl>  <dbl>    <dbl>
 1 MAZDA   MAZDA3       1        0    2006          3  89046 21973  33619     7100
 2 DODGE   1500 R...    2        0    2004          5  93593 19638  33619     7600
 3 DODGE   STRATU...    3        0    2005          4  73807 19638  33619     4900
 4 DODGE   NEON         4        0    2004          5  65617 19638  33619     4100
 5 FORD    FOCUS        5        0    2005          4  69367 19638  33619     4000
 6 MITSU... GALANT...   6        0    2004          5  81054 19638  33619     5600
 7 KIA     SPECTRA      7        0    2004          5  65328 19638  33619     4200
 8 FORD    TAURUS       8        0    2005          4  65805 19638  33619     4500
 9 KIA     SPECTRA      9        0    2007          2  49921 21973  33619     5600
10 FORD    FIVE H...   10        0    2007          2  84872 21973  33619     7700
# ... with 72,973 more rows, and 3 more variables: IsOnlineSale <dbl>,
#    WarrantyCost <dbl>, Odo_2_y <dbl>
```

Applying functions with `across` from `dplyr`

`across()` makes it easy to apply the same transformation to multiple columns, allowing you to use select() semantics inside functions like `summarize()` and `mutate()`.

`across( .cols = <columns>, .fns = function, ... )`

- List columns first : `.cols =`
- List function next: `.fns =`
- Then list any arguments for the function

Applying functions with `across` from `dplyr`.

Combining with `summarize()`:

```
cars_dbl %>%
  group_by(Make) %>%
  summarize(across(.cols = everything(), .fns = mean))
```

```
# A tibble: 33 x 12
   Make    Model  RefId IsBadBuy VehYear VehicleAge VehOdo  BYRNO VNZIP1 VehBCost
   <chr>   <dbl>  <dbl>    <dbl>   <dbl>      <dbl>  <dbl>  <dbl>  <dbl>    <dbl>
 1 ACURA      NA 36021.   0.273   2003.       6.52 81732. 21851. 61217.    9039.
 2 BUICK      NA 35431.   0.157   2004.       5.65 76238. 19755. 51298.    6169.
 3 CADIL…     NA 34173.   0.152   2004.       5.24 73770. 20383. 50775.   10958.
 4 CHEVR…     NA 35417.   0.0975  2006.       3.97 73390. 26912. 58874.    6835.
 5 CHRYS…     NA 37614.   0.129   2006.       3.65 66814. 31268. 58562.    6507.
 6 DODGE      NA 36851.   0.103   2006.       3.75 68261. 36094. 58788.    7047.
 7 FORD       NA 36866.   0.154   2005.       4.75 76749. 19887. 59427.    6403.
 8 GMC        NA 35245.   0.116   2004.       5.61 79273. 18802. 58113.    8342.
 9 HONDA      NA 35109.   0.109   2004.       5.33 77877. 24161. 52659.    8350.
10 HUMMER     NA 19533   0        2006        3    70809  21053  95673    11920
# … with 23 more rows, and 2 more variables: IsOnlineSale <dbl>,
#   WarrantyCost <dbl>
```

Applying functions with `across` from `dplyr`.

Adding arguments to the function (`quantile()`) at the end:

```
cars_dbl %>%
  group_by(Make) %>%
  summarize(across(.cols = where(is.double), .fns = quantile, probs = 0.95))
```

```
# A tibble: 33 x 11
   Make         RefId IsBadBuy VehYear VehicleAge VehOdo  BYRNO VNZIP1 VehBCost
   <chr>        <dbl>    <dbl>   <dbl>      <dbl>  <dbl>  <dbl>  <dbl>    <dbl>
 1 ACURA       67522.        1    2005       8    93338. 36099. 92807    12093
 2 BUICK       67803.        1    2007       8.05 95049. 52117  92337     8345.
 3 CADILLAC    68611.        1    2006       7    87267. 34482. 85115.   11094
 4 CHEVROLET   68895.        1    2008       7    92505  99750  94544     9170
 5 CHRYSLER    69029.        1    2008       7    89784. 99761  92504     9280
 6 DODGE       68446.        1    2008       7    91557. 99761  92504    10265
 7 FORD        69731.        1    2007       8    95213. 52117  92807     9834
 8 GMC         69012.        1    2006       8    94470  25100  92504    10912
 9 HONDA       69827         1    2007       8    93811  99740  92504    10440
10 HUMMER      19533         0    2006       3    70809  21053  95673    11920
# … with 23 more rows, and 2 more variables: IsOnlineSale <dbl>,
#   WarrantyCost <dbl>
```

Applying functions with `across` from `dplyr`.

Using different `tidyselect()` options:

```
cars_dbl %>%
  group_by(Make) %>%
  summarize(across(.cols = starts_with("Veh"), .fns = mean))
```

```
# A tibble: 33 x 5
   Make        VehYear VehicleAge VehOdo VehBCost
   <chr>         <dbl>      <dbl>  <dbl>    <dbl>
 1 ACURA         2003.       6.52 81732.    9039.
 2 BUICK         2004.       5.65 76238.    6169.
 3 CADILLAC      2004.       5.24 73770.   10958.
 4 CHEVROLET     2006.       3.97 73390.    6835.
 5 CHRYSLER      2006.       3.65 66814.    6507.
 6 DODGE         2006.       3.75 68261.    7047.
 7 FORD          2005.       4.75 76749.    6403.
 8 GMC           2004.       5.61 79273.    8342.
 9 HONDA         2004.       5.33 77877.    8350.
10 HUMMER        2006        3    70809    11920
# … with 23 more rows
```

Applying functions with `across` from `dplyr`.

Combining with `mutate()`:

```r
cars_dbl %>%
  mutate(across(.cols = starts_with("Veh"), .fns = round, digits = -3))
```

```
# A tibble: 72,983 x 12
    Make    Model    RefId IsBadBuy VehYear VehicleAge VehOdo BYRNO VNZIP1 VehBCost
    <chr>   <chr>    <dbl>    <dbl>   <dbl>      <dbl>  <dbl> <dbl>  <dbl>    <dbl>
 1 MAZDA   MAZDA3       1        0    2000          0  89000 21973  33619     7000
 2 DODGE   1500 R…      2        0    2000          0  94000 19638  33619     8000
 3 DODGE   STRATU…      3        0    2000          0  74000 19638  33619     5000
 4 DODGE   NEON         4        0    2000          0  66000 19638  33619     4000
 5 FORD    FOCUS        5        0    2000          0  69000 19638  33619     4000
 6 MITSU…  GALANT…      6        0    2000          0  81000 19638  33619     6000
 7 KIA     SPECTRA      7        0    2000          0  65000 19638  33619     4000
 8 FORD    TAURUS       8        0    2000          0  66000 19638  33619     4000
 9 KIA     SPECTRA      9        0    2000          0  50000 21973  33619     6000
10 FORD    FIVE H…     10        0    2000          0  85000 21973  33619     8000
# … with 72,973 more rows, and 2 more variables: IsOnlineSale <dbl>,
#   WarrantyCost <dbl>
```

Applying functions with `across` from `dplyr`.

Combining with `mutate()`:

```r
cars_dbl %>%
  mutate(across(.cols = everything(),
               .fns = str_replace_all,
               pattern = "A",
               replacement = "a"))
```

```
# A tibble: 72,983 x 12
   Make    Model    RefId IsBadBuy VehYear VehicleAge VehOdo BYRNO VNZIP1 VehBCost
   <chr>   <chr>    <chr> <chr>    <chr>   <chr>      <chr>  <chr> <chr>  <chr>
 1 MaZDa   MaZDa3   1     0        2006    3          89046  21973 33619  7100
 2 DODGE   1500 R…  2     0        2004    5          93593  19638 33619  7600
 3 DODGE   STRaTU…  3     0        2005    4          73807  19638 33619  4900
 4 DODGE   NEON     4     0        2004    5          65617  19638 33619  4100
 5 FORD    FOCUS    5     0        2005    4          69367  19638 33619  4000
 6 MITSU…  GaLaNT…  6     0        2004    5          81054  19638 33619  5600
 7 KIa     SPECTRa  7     0        2004    5          65328  19638 33619  4200
 8 FORD    TaURUS   8     0        2005    4          65805  19638 33619  4500
 9 KIa     SPECTRa  9     0        2007    2          49921  21973 33619  5600
10 FORD    FIVE H…  10    0        2007    2          84872  21973 33619  7700
# … with 72,973 more rows, and 2 more variables: IsOnlineSale <chr>,
#   WarrantyCost <chr>
```

Applying functions with `across` from `dplyr`.

Combining with `mutate()`:

```r
# Child mortality data
mort <- read_mortality() %>% rename(country = `...1`)


mort %>%
  select(country, starts_with("194")) %>%
  mutate(across(.cols = c(`1943`, `1944`, `1945`),
                .fns = replace_na,
                replace = 0))
```

```
# A tibble: 197 x 11
    country `1940` `1941` `1942` `1943` `1944` `1945` `1946` `1947` `1948` `1949`
    <chr>    <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>
 1 Afghan… NA      NA     NA      0      0      0     NA     NA     NA     NA
 2 Albania  1.53   1.31   1.48   1.46   1.43   1.40   1.37   1.41   1.37   1.34
 3 Algeria NA      NA     NA      0      0      0     NA     NA     NA     NA
 4 Angola   4.46   4.46   4.46   4.34   4.34   4.34   4.33   4.22   4.22   4.21
 5 Argent…  0.641  0.603  0.602  0.558  0.551  0.510  0.503  0.496  0.494  0.492
 6 Armenia NA      NA     NA      0      0      0     NA     NA     NA     NA
 7 Aruba   NA      NA     NA      0      0      0     NA     NA     NA     NA
 8 Austra…  0.263  0.275  0.276  0.299  0.260  0.271  0.295  0.279  0.271  0.271
 9 Austria  0.504  0.474  0.417  0.389  0.360  0.311  0.311  0.312  0.274  0.274
10 Azerba… NA      NA     NA      0      0      0     NA     NA     NA     NA
# … with 187 more rows
```

Website

Website