

Intro to R

Factors

Factors

A factor is a special character vector where the elements have pre-defined groups or 'levels'. You can think of these as qualitative or categorical variables:

```
x <- c("yellow", "red", "red", "blue", "yellow", "blue")  
class(x)
```

```
## [1] "character"
```

```
x_fact <- factor(x) # factor() is a function  
class(x_fact)
```

```
## [1] "factor"
```

Factors

Factors have **levels** (character types do not).

```
x
## [1] "yellow" "red"    "red"    "blue"   "yellow" "blue"

x_fact
## [1] yellow red    red    blue   yellow blue
## Levels: blue red yellow
```

Note that levels are, by default, in **alphanumerical** order.

Factors

Extract the levels of a factor vector using `levels()`:

```
levels(x_fact)
```

```
## [1] "blue" "red" "yellow"
```

forcats package

A package called forcats is really helpful for working with factors.



factor() vs as_factor()

`factor()` is from base R and `as_factor()` is from `forcats`

Both can change a variable to be of class factor.

- `factor()` will order **alphabetically** unless told otherwise.
- `as_factor()` will order by **first appearance** unless told otherwise.

If you are assigning your levels manually either function is fine!

as_factor() function

```
x <- c("yellow", "red", "red", "blue", "yellow", "blue")  
x_fact_2 <- as_factor(x)  
x_fact_2
```

```
## [1] yellow red    red    blue   yellow blue  
## Levels: yellow red blue
```

```
# Compare to factor() method:  
x_fact
```

```
## [1] yellow red    red    blue   yellow blue  
## Levels: blue red yellow
```

A Factor Example

We will use data on student dropouts from the State of California during the 2016-2017 school year. More on this data can be found here: <https://www.cde.ca.gov/ds/ad/filesdropouts.asp>

To preserve school anonymity, "CDS_CODE" is used in place of the individual school's name.

You can download the data from the JHU website here:

http://jhudatascience.org/intro_to_r/data/dropouts.txt

```
dropouts <- read_delim("http://jhudatascience.org/intro_to_r/data/dropouts.txt", delim = "\t")
```

```
## Rows: 59599 Columns: 20
## — Column specification —————
## Delimiter: "\t"
## chr (2): CDS_CODE, GENDER
## dbl (18): ETHNIC, E7, E8, E9, E10, E11, E12, EUS, ETOT, D7, D8, D9, D10, D11...
##
## [ Use `spec()` to retrieve the full column specification for this data.
## [ Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

dropouts

```
## # A tibble: 59,599 × 20
##   CDS_CODE ETHNIC GENDER  E7    E8    E9    E10    E11    E12    EUS    ETOT    D7
##   <chr>      <dbl> <chr>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 01100170...      1 M        0      0      0      0      1      1      0      2      0
## 2 01100170...      1 F        0      0      1      0      2      0      0      3      0
## 3 01100170...      2 M        0      0      0      0      0      1      0      1      0
## 4 01100170...      2 F        0      0      2      2      2      1      0      7      0
## 5 01100170...      3 M        0      0      0      1      0      0      0      1      0
## 6 01100170...      3 F        0      0      1      1      2      0      0      4      0
## 7 01100170...      4 M        0      0      0      1      0      0      0      1      0
## 8 01100170...      5 M        0      0     31     32     17     22      0    102      0
## 9 01100170...      5 F        0      0     26     34     30     20      0    110      0
## 10 01100170...      6 M        0      0     19     20     17     13      0     69      0
```

```
## # with 59,589 more rows and 8 more variables: D8 <dbl>, D9 <dbl>, D10 <dbl>...
```


Preparing the data

Aggregate (sum) across ethnicity and gender:

```
dropouts <-  
  dropouts %>%  
  group_by(CDS_CODE) %>%  
  summarize(  
    Freshman = sum(D9),  
    Sophomore = sum(D10),  
    Junior = sum(D11),  
    Senior = sum(D12)  
  )  
dropouts
```

```
## # A tibble: 5,507 × 5  
##   CDS_CODE Freshman Sophomore Junior Senior  
##   <chr>      <dbl>      <dbl> <dbl> <dbl>  
## 1 01100170112607      1          1      0      1  
## 2 01100170123968      0          0      0      0  
## 3 01100170130401      3          5     12     24  
## 4 01100170130419      1          2     13     36  
## 5 01100170131581      0          0      0      0  
## 6 01100176002000      0          0      0      0  
## 7 01316090131755      0          0      0      0  
## 8 01316170131763      0          0      0      2  
## 9 016111900000001      0          1      0      1  
## 10 01611190106401      0          0      0      0  
## # ... with 5,497 more rows
```

Preparing the data

Pivot to long format:

```
dropouts <-  
  dropouts %>%  
  pivot_longer(  
    !CDS_CODE,  
    names_to = "grade",  
    values_to = "n_dropouts"  
  )  
dropouts
```

```
## # A tibble: 22,028 × 3  
##   CDS_CODE      grade n_dropouts  
##   <chr>      <chr>      <dbl>  
## 1 01100170112607 Freshman      1  
## 2 01100170112607 Sophomore     1  
## 3 01100170112607 Junior        0  
## 4 01100170112607 Senior        1  
## 5 01100170123968 Freshman      0  
## 6 01100170123968 Sophomore     0  
## 7 01100170123968 Junior        0  
## 8 01100170123968 Senior        0  
## 9 01100170130401 Freshman      3  
## 10 01100170130401 Sophomore     5  
## # ... with 22,018 more rows
```

The data

```
head(dropouts)
```

```
## # A tibble: 6 × 3
##   CDS_CODE      grade  n_dropouts
##   <chr>      <chr>      <dbl>
## 1 01100170112607 Freshman      1
## 2 01100170112607 Sophomore     1
## 3 01100170112607 Junior        0
## 4 01100170112607 Senior        1
## 5 01100170123968 Freshman      0
## 6 01100170123968 Sophomore     0
```

Notice that `grade` is a `chr` variable. This indicates that the values are **character** strings.

R does not realize that there is any order related to the `grade` values. It will assume that it is **alphabetical**.

However, we know that the order is: **freshman, sophomore, junior, senior**.

For the next steps, let's take a subset of data.


Use `set.seed()` to take the same random sample each time.

```
set.seed(123)  
dropouts_subset <- slice_sample(dropouts, n = 32)
```

Plot the data

Let's make a plot first.

```
dropouts_subset %>%  
  ggplot(mapping = aes(x = grade, y = n_dropouts)) +  
  geom_boxplot() +  
  theme_bw(base_size = 16) # make all labels size 16
```



OK this is very useful, but it is a bit difficult to read. We expect the values to be plotted by the order that we know, not by alphabetical order.

Change to factor

Currently `grade` is class character but let's change that to class factor which allows us to specify the levels or order of the values.

```
dropouts_fct <-  
  dropouts_subset %>%  
  mutate(grade = factor(grade,  
    levels = c("Freshman", "Sophomore", "Junior", "Senior")  
  ))  
  
dropouts_fct %>%  
  pull(grade) %>%  
  levels()  
  
## [1] "Freshman" "Sophomore" "Junior" "Senior"
```

Change to a factor

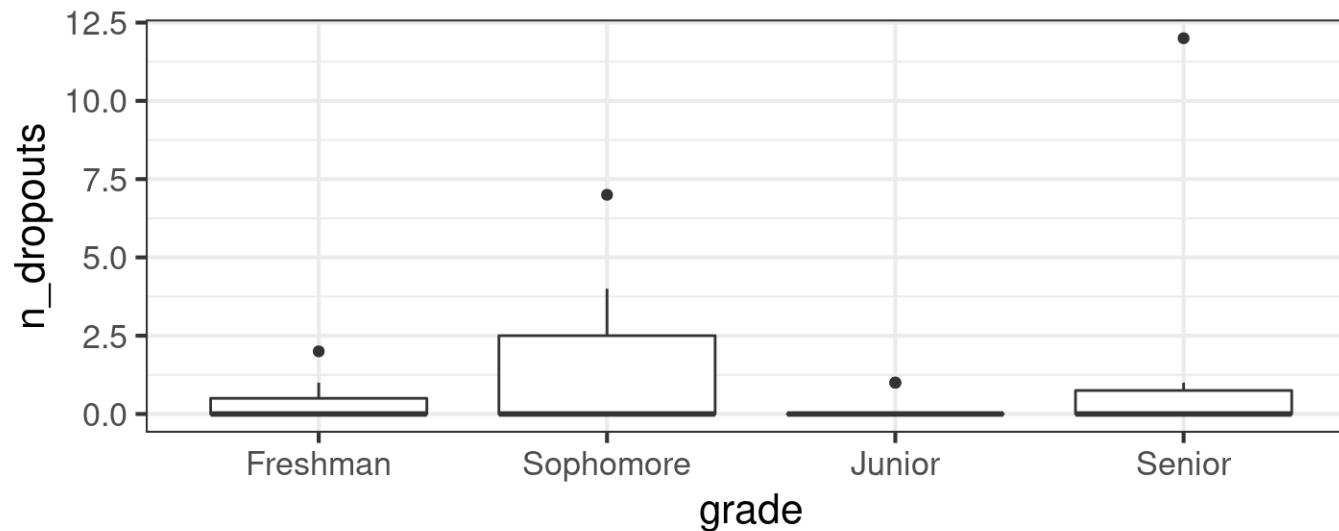
```
head(dropouts_fct)
```

```
## # A tibble: 6 × 3
##   CDS_CODE      grade  n_dropouts
##   <chr>      <fct>      <dbl>
## 1 45699716050231 Junior          0
## 2 45700036050330 Junior          0
## 3 12630401230069 Sophomore       0
## 4 09100900930131 Sophomore       1
## 5 15633216009179 Junior          0
## 6 33670330113647 Sophomore       0
```

Plot again

Now let's make our plot again:

```
dropouts_fct %>%  
  ggplot(mapping = aes(x = grade, y = n_dropouts)) +  
  geom_boxplot() +  
  theme_bw(base_size = 16)
```



Now that's more like it! Notice how the data is automatically plotted in the order we would like.

What about if we `arrange()` the data by grade?

Character data is arranged alphabetically.

```
dropouts_subset %>%  
  arrange(grade)
```

```
## # A tibble: 32 × 3  
##   CDS_CODE      grade  n_dropouts  
##   <chr>      <chr>      <dbl>  
## 1 19643941931823 Freshman          1  
## 2 11626616007611 Freshman          0  
## 3 23656150128280 Freshman          0  
## 4 30664313030616 Freshman          0  
## 5 54719935432414 Freshman          2  
## 6 22655326025050 Freshman          0  
## 7 37683386039911 Freshman          0  
## 8 45699716050231 Junior            0  
## 9 45700036050330 Junior            0  
## 10 15633216009179 Junior            0  
## # ... with 22 more rows
```

Notice that the order is not what we would hope for!

Arranging Factors

Factor data is arranged by level.

```
dropouts_fct %>%  
  arrange(grade)
```

```
## # A tibble: 32 × 3  
##   CDS_CODE      grade  n_dropouts  
##   <chr>      <fct>      <dbl>  
## 1 19643941931823 Freshman        1  
## 2 11626616007611 Freshman        0  
## 3 23656150128280 Freshman        0  
## 4 30664313030616 Freshman        0  
## 5 54719935432414 Freshman        2  
## 6 22655326025050 Freshman        0  
## 7 37683386039911 Freshman        0  
## 8 12630401230069 Sophomore        0  
## 9 09100900930131 Sophomore        1  
## 10 33670330113647 Sophomore        0  
## # ... with 22 more rows
```

Nice! Now this is what we would want!

Making tables with characters

Tables grouped by a character are arranged alphabetically.

```
dropouts_subset %>%  
  group_by(grade) %>%  
  summarize(total_dropouts = sum(n_dropouts))
```

```
## # A tibble: 4 × 2  
##   grade      total_dropouts  
##   <chr>          <dbl>  
## 1 Freshman           3  
## 2 Junior             2  
## 3 Senior            13  
## 4 Sophomore         12
```

Making tables with factors

Tables grouped by a factor are arranged by level.

```
dropouts_fct %>%  
  group_by(grade) %>%  
  summarize(total_dropouts = sum(n_dropouts))
```

```
## # A tibble: 4 × 2  
##   grade      total_dropouts  
##   <fct>          <dbl>  
## 1 Freshman           3  
## 2 Sophomore        12  
## 3 Junior            2  
## 4 Senior          13
```

forcats for ordering

What if we wanted to order `grade` by increasing `n_dropouts`?

```
library(forcats)

dropouts_fct %>%
  ggplot(mapping = aes(x = grade, y = n_dropouts)) +
  geom_boxplot() +
  theme_bw(base_size = 16)
```



This would be useful for identifying easily which grade to focus on.

forcats for ordering

We can order a factor by another variable by using the `fct_reorder()` function of the `forcats` package.

```
library(forcats)

dropouts_fct %>%
  ggplot(mapping = aes(x = fct_reorder(grade, n_dropouts, mean), y = n_dropouts)) +
  geom_boxplot() +
  labs(x = "Student Grade") +
  theme_bw(base_size = 16)
```



Adding another variable

Let's say that we also want to assess which grade has the most incidences of being tardy (another word for late) to class. - Now we will add another simulated variable of random values from 0 to 7 and of 32 values total. - We set a seed again so that our results will be consistent each time we run this code.

```
set.seed(1956)
dropouts_fct <-
  dropouts_fct %>%
  mutate("tardy" = sample(0:7, size = 32, replace = TRUE))

head(dropouts_fct)
```

```
## # A tibble: 6 × 4
##   CDS_CODE      grade  n_dropouts  tardy
##   <chr>      <fct>      <dbl> <int>
## 1 45699716050231 Junior         0         1
## 2 45700036050330 Junior         0         3
## 3 12630401230069 Sophomore      0         5
## 4 09100900930131 Sophomore      1         1
## 5 15633216009179 Junior         0         4
## 6 33670330113647 Sophomore      0         0
```

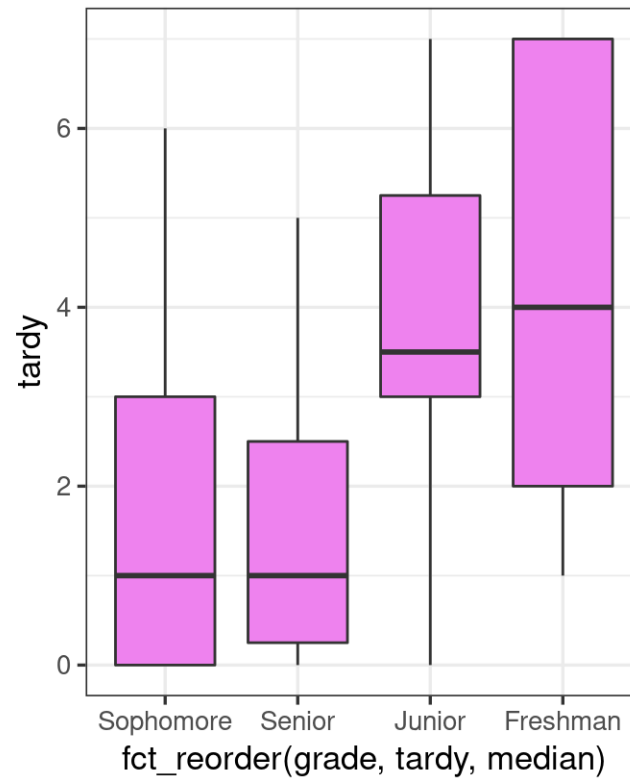
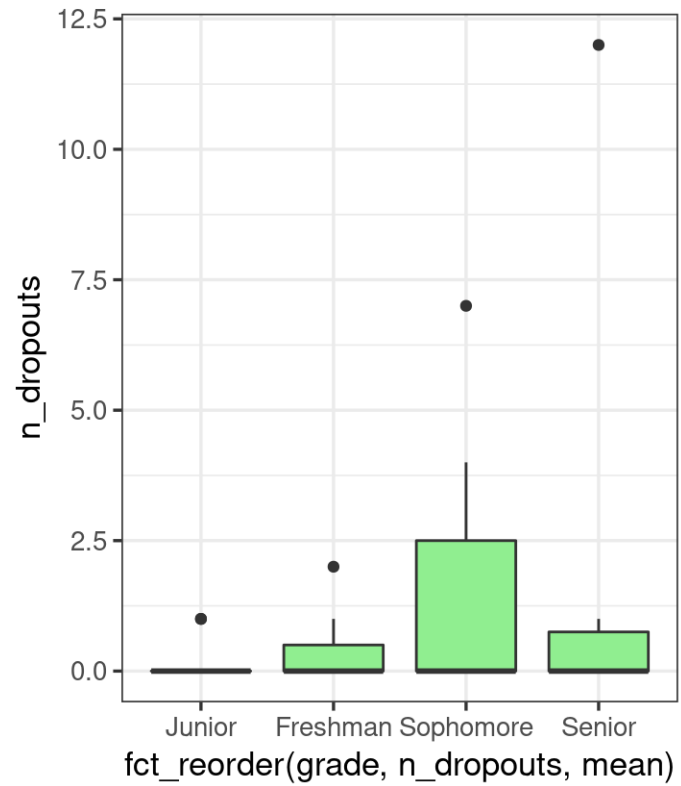
Plotting new variable

Now let's plot each of our variables of interest (n_dropouts and tardy) on the y axis and grade on the x axis. Let's arrange grade by the amount of each.

```
dropouts_plot <- dropouts_fct %>%  
  ggplot(mapping = aes(  
    x = fct_reorder(grade, n_dropouts, mean),  
    y = n_dropouts  
  )) +  
  geom_boxplot(fill = "lightgreen") +  
  theme_bw(base_size = 13)  
  
tardy_plot <- dropouts_fct %>%  
  ggplot(mapping = aes(  
    x = fct_reorder(grade, tardy, median),  
    y = tardy  
  )) +  
  geom_boxplot(fill = "violet") +  
  theme_bw(base_size = 13)  
  
library(patchwork)  
dropouts_plot + tardy_plot
```



Plotting new variable

The result!



Cleaning up the plots

```
dropouts_plot +  
  labs(x = "Student Grade", y = "Number of Dropouts") +  
  tardy_plot +  
  labs(x = "Student Grade", y = "Tardy Days")
```



fct_count

The `fct_count()` function of the `forcats` package is helpful for checking that the proportions of each level for a factor are similar. Need the `prop = TRUE` argument otherwise just counts are reported.

```
dropouts_fct %>%  
  pull(grade) %>%  
  fct_count(prop = TRUE)
```

```
## # A tibble: 4 × 3  
##   f           n     p  
##   <fct>     <int> <dbl>  
## 1 Freshman      7 0.219  
## 2 Sophomore     7 0.219  
## 3 Junior      12 0.375  
## 4 Senior        6 0.188
```

Summary

- the factor class allows us to have a different order from alphanumeric for categorical data
- we can change data to be a factor variable using `mutate` and a factor creating function like `factor()` or `as_factor`
- the `as_factor()` is from the `forcats` package (first appearance order by default)
- the `factor()` base R function (alphabetical order by default)
- with `factor()` we can specify the levels with the `levels` argument if we want a specific order
- the `fct_reorder({variable_to_reorder}, {variable_to_order_by}, {summary function})` helps us reorder a variable by the values of another variable
- arranging, tabulating, and plotting the data will reflect the new order

Lab

[Class Website](#)
[Lab](#)



Image by [Gerd Altmann](#) from [Pixabay](#)